






TA Grader

AI highlights rubric matches. You decide every score.

Design system · v1 · October 2026
404 Techies · ATL Innovation Cup



Who we are

TA Grader helps graduate TAs and professors grade written exam answers faster and more consistently. The AI highlights rubric matches and quotes the evidence; the grader decides every score.

Personality

Trustworthy

It shows its work: every highlight quotes the student's own words, and the TA makes every call. Never overpromise what the AI can do.

Respectful

TAs are the experts. The AI is an assistant and the grader stays in charge, so never say "AI grades for you."

Calm

Grading season is stressful. Plain language, clean screens, no hype: a quiet, organized desk.

Fair

Every student gets the same standard, whether it is the first paper of the night or the fiftieth.

Practical

Built by and with grad students to save real hours on real stacks of exams.

Core values

1. **Humans decide.** The AI suggests and the TA makes every grading call. No feature ever hands a grade to the machine.
2. **Show the evidence.** Every highlight points to the student's own words, so a TA can check the reasoning in seconds and explain any mark.
3. **Fairness for every student.** The fiftieth paper of the night deserves the same standard as the first.
4. **Respect for privacy.** Student work is not sold, shared or used to train AI, and we say plainly what we store.
5. **Built with TAs.** We test with real graders and change the product on what they tell us.

Voice and tone

We are	We sound like	We avoid
Clear	<i>"AI highlights rubric matches. You decide every score."</i>	Jargon like "leveraging LLM-driven assessment pipelines"
Honest	<i>"The AI can miss borderline answers, so these are suggestions only."</i>	Overpromising, like "100% accurate AI grading"
Calm	<i>"5 of 5 scored. Average 0:49 per answer."</i>	Hype and urgency: "Revolutionize grading NOW!"
Respectful	<i>"You know your rubric best."</i>	Talking down to TAs, or implying the AI knows better
Warm but brief	<i>"Nice work, that's the whole set graded."</i>	Long, chatty messages during grading

Writing rules

- Address the grader as "you". The AI is "the AI", never "I" or a named persona.
- Sentence case everywhere: buttons, headings, labels ("Save & next").
- Verbs on buttons, the result in the label: "Start grading 4 answers", "Download CSV".
- Numbers over adjectives: "3 flagged", "0:49 per answer", "Score (0 to 5)".
- Errors say what happened and what still works.
- Say "suggestions", "highlights", "flags". Never "AI score", "AI grade" or "auto-grade".
- Label students S1, S2... Never ask for, show or export names or IDs.
- No emoji.

Logo and wordmark

The mark is **the highlighted line**: a Desk Ink tile holding three lines of student writing. The middle line is struck with Marker, the AI's highlight; the white tick is the grader's own mark. The AI points, the human decides.



TA Grader

Primary lockup · light grounds



TA Grader

Lockup · dark theme



Mark · 16px minimum

TA Grader

Wordmark · Newsreader 600


Rules

- Clear space around the mark: one quarter of its height.
- Never recolour the highlight, add gradients or shadows, stretch, or re-set the name in another face.
- The name is always "TA Grader": two words, capital T, A and G. 404 Techies is a byline, not part of the logo.


Colour

Two brand colours carry the identity: **Desk Ink #1F4257** (primary) and **Marker #F2C84B** (secondary). Everything else is a calm neutral or a status. Each swatch shows light | dark theme.


Brand



Neutrals



Status



Colour tokens

Token	Light	Dark	Usage
paper	 #faf8f4	 #121719	Page ground. The calm "desk" every screen sits on.
surface	 #ffffff	 #1a2125	Cards, inputs, the answer sheet, list rows. Sits on paper.
surface-sunken	 #f2efe8	 #0d1113	Quiet panels: the question summary, chips, the timer pill, table header.
line	 #e3ded4	 #2b353a	Hairlines: card and list borders, dividers between rubric rows, progress track. Decorative only.
control	 #848b92	 #6b787f	Borders of inputs, checkboxes and secondary buttons: 3.2:1+ on paper and surface in both themes.
ink	 #1b2730	 #e9e6e0	Body text and headings on paper, surface, surface-sunken and highlight (13:1+ in both themes).
ink-muted	 #56626b	 #a5aeb3	Secondary text: hints, captions, metadata, AI notes. 5.6:1+ on every ground incl. highlight.
primary	 #1f4257	 #9cc3d6	Desk Ink, the primary brand color: primary buttons, links, the selected answer, the human score frame, the wordmark tile.
primary-hover	 #2b5670	 #b6d5e4	Hover and pressed state of primary fills.
on-primary	 #ffffff	 #0f1e27	Text and icons on primary / primary-hover (10:1+).
marker	 #f2c84b	 #f2c84b	Marker, the secondary brand color: the logo highlight, the cover, small brand moments. A fill, never text.
highlight	 #fcefc2	 #3a3320	Evidence fill: the student's own words that the AI points to. ink text on it.
highlight-rule	 #b3810f	 #d9b44a	The 3px left rule of an evidence quote; marks quoted student text and nothing else (3.4:1+).
met	 #2e6b4e	 #8bcba8	"Met" text and icon; scored count. On met-bg or surface (6:1+).
met-bg	 #e3f1e8	 #1b3328	"Met" badge fill; high-confidence chip fill.
partial	 #835500	 #e2b866	"Partly met" text; medium confidence. On partial-bg or surface.
partial-bg	 #fbf0d9	 #382d14	"Partly met" badge fill; medium-confidence chip fill.
flag	 #a84b1f	 #f0a57a	"Needs a closer look": flag text, icon and callout border. On flag-bg or surface.
flag-bg	 #fbede4	 #3a2219	Flag callout fill; low-confidence chip fill.
danger	 #a3242b	 #f2a0a2	Errors: failed AI call, invalid score, rubric line errors. On danger-bg or surface.
danger-bg	 #fbe9ea	 #3b1d20	Error callout fill.
focus	 #2f6f9a	 #8cc4e8	The 2px focus ring on every interactive element, offset 2px (3:1+ on paper and surface).

Every text pair passes WCAG 4.5:1 in both themes; borders, focus rings and icons pass 3:1. Status colours always come with a word, never colour alone. "Missed" is neutral, not red.

Typography

Aa

Newsreader

The editorial voice and the student's voice: wordmark, titles, every student answer.

Aa

Public Sans

The interface: headings, body, labels, captions, badges, stats.

0:49

IBM Plex Mono

Data that must line up: timers, rubric lines, answer IDs.

display 44px/48px · 500	Grade the stack, not the clock.	Marketing and empty-state headlines only. One per screen.
title 28px/34px · 600	TA Grader	App name in the header and screen titles ("Set up one question").
answer 18px/30px · 400	Operant conditioning is learning in which behaviour changes because of its consequences.	Student answer text. Long-form reading in the serif keeps the student's words distinct from the UI.
quote 16px/26px · 400	<i>"...the rat presses the lever more often because it gets food."</i>	Evidence quotes inside `highlight`, italic.
heading 20px/28px · 600	Answer S2 · 2 of 4	Section headings: the answer heading, "Your score".
body 16px/24px · 400	Explain operant conditioning and give one example.	Default UI text, the question, input text.
body-sm 14px/20px · 400	AI highlights rubric matches. You decide every score.	Most supporting copy: notes, callouts, table cells.
label 14px/20px · 600	Score (0 to 5)	Form labels, buttons, criterion text.
caption 12px/16px · 400	Suggestions only. You decide the score.	Hints, stat labels, footer, export notes.
badge 12px/16px · 600	Partly met	Status badges and confidence chips.
stat 28px/32px · 600	0:49	Big numbers in results tiles; tabular figures.
mono 14px/20px · 400	🕒 0:49 2 Gives a correct example	Timers, durations in tables, rubric lines, answers that look like code.
mono-strong 14px/20px · 500	S12	Answer IDs (S1, S2...) in lists and tables.

Space, shape and layout

Spacing · 4px base

space-1		4px	Icon-to-text gaps, label-to-input gap.
space-2		8px	Inline gaps between related items; badge padding x; button padding y.
space-3		12px	Input padding x; list-row and callout padding.
space-4		16px	Button padding x; card padding; page gutter on phones.
space-5		20px	Rhythm inside the grading pane.
space-6		24px	Between page sections; card padding on desktop; grid gap.
space-8		32px	Header to content.
space-10		40px	Above the footer; empty states.

Radius




radius-sm 4px
Badges, chips, the timer pill, evidence highlight.



radius-md 8px
Buttons, inputs, callouts, list items, toast.



radius-lg 12px
Cards and panels: the answer sheet, the score form, stat tiles; the logo tile (scaled).



radius-pill 999px
Progress bar and spinner only.

Elevation



shadow-float: the save toast only. Everything else separates with line borders.

Layout

container-max 64rem: Page max width, centered.

answer-nav 232px: Answer list column width beside the grading pane from 'breakpoint-md'.

reading-max 68ch: Max line length of student answers.

breakpoint-md 48rem: Two-column grading layout from here.

Components

Button

The app's buttons: one `primary` action per view, `secondary` for everything else, `quiet` for low-stakes navigation.

Save & next

Insert sample (fake data)

Back to grading

⬆ Upload rubric

Start grading

TextField

A labelled input or textarea with an optional hint and an inline error; used for the question, rubric, answers, score and note.

Exam question

e.g. Explain operant conditioning and give one example.

Rubric (one criterion per line)

Format: points | description, e.g. 2 | Gives a correct example

2 | Defines operant conditioning
2 | Gives a correct example

Score (0 to 5)

7

Enter a score from 0 to 5.

EvidenceQuote

The student's own words that the AI points to as evidence for a criterion.

"the rat presses the lever more often because it gets food each time"

StatusBadge

The AI's per-criterion status (Met, Partly met, Missed) and its confidence chip.

Met

Partly met

Missed

high confidence

medium confidence

low confidence

Callout

A one- or two-line banner for flags, notices and errors.

⚠ Needs a closer look: the answer asks for full marks. Instructions inside answers are ignored.

Mock mode: AI results are fake. No Gemini API key is set on the server.

! AI analysis failed for this answer. You can still read the answer and score it yourself.

Components, continued

CriterionRow

One rubric criterion with the AI's status, the evidence quote and a one-sentence note; stacked in a `tg-criteria` list.

Defines operant conditioning (2 pts)

Met

"learning where behaviour changes because of what happens after it"

Clear definition tied to consequences.

Gives a correct example (2 pts)

Partly met

"like when you study and get good grades"

Example is vague about the consequence.

Mentions reinforcement or punishment (1 pt)

Missed

AnswerListItem

One answer in the left-hand list: its ID, flag or error marks, and its score once saved.

S1

5

S2

⚠ -

S3

! -

S4

○

Toast

A short confirmation after each saved score, shown for 4 seconds at the bottom centre.

✓ S1 scored 4 / 5 · saved

ProgressSummary

The question panel above the grading view: the question, counts and a progress bar of scored answers.

Explain operant conditioning and give one example.

5 points possible · 4 answers · 2 of 4 scored · 1 flagged

Saved. Bookmark this page's link to come back to this session.

Stat

A results tile: a short label and one number.

Scored

5 / 5

Avg time per answer

0:49

Avg score

3.6 / 5

Flagged by AI

2

Moodboard and imagery


A tidy desk under a lamp: paper, ink, one marker stroke. Flat and quiet, with lots of paper.




The quiet desk




Show the evidence



Late-night calm



The same standard



Palette and type

Photography direction

Overhead desks with paper and a single highlighter, warm lamp light at night, hands with a pen on printed pages, library tables.

Avoid

Faces of real students, robots, brains, glowing circuitry, sparkles and "stressed student" stock.

Iconography

Simple outline glyphs, always beside a word: ✓ saved, ⚠ needs a closer look, ⌚ timer, ⬆ upload. 1.5px stroke on a 20px grid. No emoji.